AI-based NPS improvement

> For reliance retail which is physical sores and online Jio mart..ecomm and quick comm shopping and delivery service.. I want to create a next generation reqmts document which measures customer store footfall and experience nps and also digital eomm /q comm service nps without asking the customer any questions and survey but in real time signal collection and analysis..i want this to be retail employee actions, policy corrections and sop upgradation so we continuously improve our nps by our actions and periodically confirm with a a third party customer survey so we consistently deliver 10 x higher nps than our competitors in market..pl explain form first principles the logic, use today's Ai first tech and align policy, workflows, sops to an Ai first architecture with our 12 layered foundation platform architecture so we get real time actions for org and continuous improvement.. Pl improve this prompt first and ask any questions you have to give me an understable unique design basis and rationale that is disruptive and a practical execution plan, org and recommendations on specialist partners who can accelerate my execution

Request for GPT-5 Pro
Reasoned for 6m 37s
Fantastic ambition. I've taken your idea ("measure in-store + digital NPS without surveys; turn signals into real-time staff actions; calibrate periodically with an external survey; architect everything AI-first on your 12-layer platform") and turned it into:

1. a stronger prompt you can reuse internally or with vendors,
2. a first-principles design (logic + math for "silent NPS"),
3. an AI-first, 12-layer reference architecture aligned to policies, workflows, SOPs,
4. a practical execution plan (quarters → weeks),
5. org design & operating model,
6. a short list of specialist partners (global + India) to accelerate delivery (with sources).

I've also included a tight set of clarification questions at the end. They are optional—this plan stands on its own now.

---

1) Your improved prompt (copy/paste)

Goal
Design and implement an AI-first, survey-free system that continuously infers customer satisfaction/NPS in Reliance Retail's physical stores and on JioMart (e-comm and quick-commerce), converts signals into real-time frontline actions and SOP corrections, and validates impact via periodic third-party NPS surveys. Target: materially higher (10× delta vs competitors is the aspiration) sustained NPS and cost-to-serve reductions.

Business outcomes (12 months)

- ↑ In-store conversion vs traffic; ↓ queue & service wait times; ↓ out-of-stocks; ↑ on-time delivery (OTD) & ETA accuracy; ↓ support contacts per order; ↑ repeat purchase and retention.
- Continuous "silent NPS" (inferred NPS) with actionability at store/slot/region/app-journey level; statistically verified uplift via controlled experiments.
- Policy-as-code SOPs that auto-trigger store and online actions with auditability and privacy-by-design.

Scope

- **Signals (no surveys):** in-store video/IoT/POS; app/web telemetry; search/checkout/payment signals; OMS/WMS/logistics; support chat/voice transcripts; social/app-store reviews.
- **Decisioning:** real-time triggers for staff & systems; experiment-ready; causal impact measurement.
- **Governance:** DPDP-compliant consent, minimization, retention; identity resolution with consent managers; security, audit, model governance. (DPDP establishes consent managers and fiduciary obligations. Future of Privacy Forum+1)
- **Validation:** quarterly (or semi-annual) third-party NPS benchmark.

Non-goals

- Facial recognition or personally identifying video analytics; any action that breaches DPDP or internal policy.

Deliverables

- Signal catalog & "silent NPS" spec, feature store definitions, policy-as-code library, SOP playbooks, store & ops-copilot apps, dashboards, experiment framework, and rollout plan.

Constraints & SLOs

- Hot path (signal→action): ≤5s; Warm path insights: ≤60m; Cold path: daily.
- Edge processing for video; privacy filters before cloud; opt-in identity linking.
- Reliability SLO ≥ 99.9% for actioning.

---

2) First-principles logic: "silent NPS" without surveys

What we want really: NPS = %Promoters − %Detractors (from a one-question survey). Since we won't ask, we **estimate the probability** each interaction would have been a promoter/detractor **from behavior and context**. We then calibrate this estimate against periodic external NPS.

a) Signals that imply delight/pain (illustrative)

- **Store:** entry counts, dwell by zone, product-find events, associate engagement, queue length & wait, checkout success/time, OOS & substitution rates, returns/complaints. (Mature people-counting & occupancy exist; e.g., RetailNext, Sensormatic/ShopperTrak. retailnext.net+3)
- **Digital:** journey drop-offs, error codes, performance (TTFB/LCP), "frustration signals" (rage clicks, thrashes, dead clicks), payment retries, cancellations, re-orders. (Platforms like FullStory/Quantum Metric/Contentsquare formalize frustration signals & continuous product design. fullstory.com+2)

- **Delivery:** promised vs actual ETA, freshness/damage claims, first-attempt-delivery rate (FADR), courier-customer contact, geo-stops. (All-mile platforms like Locus provide route/OTD optimization and FADR metrics. Locus+1)
- **Voice/text:** call/chat sentiment, topic, resolution, escalations (Uniphore, Observe.AI). Uniphore+1
- **Public signals:** social listening, app-store reviews (Sprinklr/Talkwalker; AppFollow). Sprinklr+2

b) Model the latent outcome

- Define label $Y \in \{-1, 0, +1\}$ (detractor, passive, promoter). Train a model to estimate $P(Y = +1)$ and $P(Y = -1)$ from features built on the signals above, using **ground-truth** from periodic, stratified NPS surveys + strong proxies (e.g., repeat purchase within X days *after a positive experience*, successful substitutions, resolved complaints).
- Convert to **silent NPS (sNPS)** per entity (store/day, slot, journey step):

$$\text{sNPS} = \frac{\sum(P_{+1} - P_{-1})}{N} \times 100$$

- Calibrate with **isotonic regression** or **Platt scaling**, then **re-calibrate** after each external NPS wave; track drift.

c) Make it causal & fair

- Use **uplift models** and **difference-in-differences** to estimate *causal* impact of actions on sNPS.
- Guardrails: exclude protected attributes; design for **counterfactual fairness**; DPDP compliance (consent, minimization, rights to erasure). (DPDP obligations for data fiduciaries and consent managers apply. MeitY+1)

d) Turn sNPS into action

- Define **thresholds & triggers**: e.g., *Queue wait > 7 min* → open counter task; *Shelf gap detected* → restock task; *Delivery ETA slip > 10 min* → proactive SMS + adjust routing; *Frustration spike on checkout step* → rollback feature flag or page fix; *Negative call sentiment + unresolved* → callback within 30 minutes. (Queue/occupancy solutions and digital frustration analytics are industry-standard. Qmatic+2)

---

3) AI-first, 12-layer foundation platform (mapped to policies, workflows, SOPs)

L1 Edge Sensing & Telemetry

- **Store:** camera analytics (people counting, dwell, queue), shelf OOS, ESLs, POS, RFID; run **on-edge** for privacy; devices like Axis (people counter) or NVIDIA Metropolis partners. Axis Communications+2
- **Digital:** app/web SDKs for events, errors, performance; OpenTelemetry for traces/metrics/logs (**standard**). OpenTelemetry
- **Logistics:** courier app events, telematics, geo.

L2 Identity, Consent & Privacy

- Consent capture, preference store, DPDP **Consent Manager** integration, pseudonymous IDs, **hash-phone** linking for loyalty; on-edge face blurring (no facial ID). (DPDP's consent manager concept is core. Future of Privacy Forum)

L3 Ingestion & Streaming

- Event bus (Kafka/Confluent) with schema registry; privacy filters before persistence. Confluent+1

L4 Real-time Processing & CEP

- Stream processing (Flink/Spark-streaming) for sessionization, joins, queue-time, ETA deviations, anomaly detection. flink.apache.org+1

L5 Lakehouse & Stores

- Batch & streaming storage (Delta/Parquet), time-series; warehouse/lakehouse on Databricks/Snowflake. Databricks+1

L6 Feature Store

- Online/offline feature parity; reusable features (e.g., "checkout_error_rate_5m", "queue_wait_p95_10m"). Options: Tecton/Feast. Tecton+1

L7 Modeling & Causal Inference

- sNPS models, uplift models, forecast OOS/queue; topic/sentiment models for CX transcripts.

L8 Online Inference & Rules/Policy

- Low-latency scoring (<50–100ms) + **policy-as-code** guardrails (OPA/Cedar). openpolicyagent.org+1

L9 Decisioning & Orchestration

- Event-driven workflows (e.g., Temporal-like or cloud equivalent) to create **tasks** in store ops apps, OMS/WMS, WFM, or customer messaging; real-time **feature flags** to roll back bad releases. (Experimentation/flags vendors below. Optimizely)

L10 Experimentation & Causal Platform

- A/Bs, CUPED, geo-experiments; guardrails; platform choices include Statsig/Eppo/Optimizely. Statsig+2

L11 Observability & Experience Analytics

- OpenTelemetry across services; DX analytics for frustration heatmaps & replay (FullStory/Contentsquare/Glassbox). OpenTelemetry+2

L12 Governance, Risk, Compliance & Catalog

- Data catalog/lineage (Atlan), model registry (MLflow), DPIA templates, RBAC/ABAC, audit, retention. Atlan+1

  Policy alignment: Encode SOPs as **policy-as-code** (L8) so changes to SOPs are versioned, testable, and auto-actioned when signals breach thresholds (e.g., "If queue_p95 > 7m for

10m, notify SM + open till + greet waiting customers”). (OPA/Cedar fit this “policy-as-code” approach. openpolicyagent.org+1)

4) Signal catalog → actions (examples you can pilot fast)

Journey	Key silent signals	Action (policy-as-code)	Owner
Entrance → Browse (Store)	Entry count, dwell variance, aisle abandonment, OOS shelf gaps	Trigger *Floor Assist* if dwell>threshold & no engagement; task *Restock* if gap with stock in backroom	Store Ops
Queue & Checkout (Store)	Queue length & wait, declined payments	*Open checkout* task to nearest trained associate; proactively guide to self-checkout; payment terminal health check	SM + Cashiers
Search → PDP → Checkout (App/Web)	Rage clicks, dead clicks, errors, slow LCP/TTFB	Auto-rollback new UI flag; file bug; show lightweight alternative checkout; notify on-call	Product/Engg
Promise→Deliver (Quick-Comm)	ETA deviation, courier dwell, substitution quality	“OTD breach risk” trigger → reroute courier or auto-notify customer + goodwill coupon within cap	Fulfillment
Support	Negative sentiment + unresolved 1st contact	Supervisor callback within 30m; embed “fix forward” step into app post-call	Care

(Frustration analytics and occupancy/people counting are established practices. fullstory.com+1)

5) Practical execution plan (no waiting for surveys)

Phase 0 (Weeks 0–2): Blueprint & guardrails

- Approve principles: DPDP compliance, “no facial ID”, consent flows; define **hot/warm/cold SLOs**; choose 3–5 KPIs for pilot (e.g., queue wait, OTD, checkout errors, sNPS). (DPDP obligations & consent manager model. MeitY+1)

Phase 1 (Weeks 2–10): Instrument & pilot (3 stores + 1 city for quick-comm + 5% of app traffic)

- Stores: install/enable people-counting & queue analytics on existing cameras (Axis/Metropolis partners or RetailNext/Sensormatic), POS events, shelf OOS in 1–2 categories. Axis Communications+2
- Digital: add OpenTelemetry + a DX analytics SDK (choose one), instrument rage clicks & errors; stand up Kafka/Flink hot path; create 10 core features in feature store. OpenTelemetry+2
- Logistics: integrate ETA/route events from Locus (or internal). Locus
- Launch **Ops Copilot** (mobile): receive tasks, acknowledge, escalate; start with 6 playbooks (queue, OOS, ETA slip, app checkout error spike, payment retry spike, negative call sentiment).

Phase 2 (Weeks 10–18): sNPS model + actions

- Train v1 sNPS using proxies (repeat, refund/complaint signals, app ratings) and a small external NPS sample for calibration; wire **action policies** for top 10 triggers in OPA/Cedar. openpolicyagent.org+1

- Stand up experimentation platform (Eppo/Optimizely/Statsig) and run **two geo or store-level A/Bs** (e.g., queue playbook; OTD notifications). geteppo.com+1

Phase 3 (Months 4–6): Scale to 100 stores + 1–2 metros; 25–30% app traffic

- Add more signals (OOS computer vision, substitution quality); expand action playbooks; start **weekly policy retro** to patch SOPs where triggers fire often. (Trax/Wobot for shelf & SOP compliance are proven categories. Trax Retail+1)

Phase 4 (Months 7–12): Nationwide scale

- Nationwide rollout; quarterly external NPS benchmark (Bain/Satmetrix/Medallia partner). (NPS method and leading providers. Bain+1)

---

6) Org design & ways of working

- **Head of Experience Intelligence (HXI)** reporting to Retail CEO; owns sNPS & experimentation.
- **Signals & Causal Lab** (data science + causal inference) → owns sNPS models, uplift analysis, calibration with survey vendor.
- **Edge AI & Store Tech** → cameras, edge compute, policy-as-code deployments.
- **Ops Copilot & Activation** → playbooks, agent apps, workforce integration.
- **Privacy & AI Governance** → DPDP compliance, model risk, audits, data catalog/lineage (Atlan), model registry (MLflow). Atlan+1
- **Experimentation Guild** → standards for A/Bs, results review, rollout/rollback.

---

7) Specialist partners (curated, pick 1–2 per lane)

In-store footfall/queue & shelf health

- **Sensormatic/ShopperTrak** — mature traffic analytics, employee exclusion for accurate counts. sensormatic.com+1
- **RetailNext** — people counting, real-time occupancy & conversion analytics. retailnext.net+1
- **Axis Communications** (hardware + on-edge analytics like People Counter; 3D counters). Axis Communications+1
- **Trax** — computer vision for shelf OOS and signal-based merchandising. Trax Retail
- **Wobot.ai** (India) — video intelligence for SOP compliance on existing CCTV. Wobot AI
- **NVIDIA Metropolis ecosystem** — partner network & tooling for vision AI at the edge. NVIDIA

Digital experience analytics (DXA)

- **FullStory / Quantum Metric / Contentsquare / Glassbox** — session replay, frustration signals, tagless analytics at scale. fullstory.com+3

All-mile optimization (quick-commerce)

- **Locus.sh** — dispatch, route optimization, FADR, OTD visibility (India-born, global). Locus

Contact-center AI & conversation intelligence

- **Uniphore / Observe.AI** — analyze 100% of calls/chats, sentiment, auto-QA, coaching. Uniphore+1

Streaming, features, ML operations

- **Confluent (Kafka)** for streaming; **Apache Flink** for real-time computation; **Tecton/Feast** for feature store; **MLflow** for model registry. Confluent+4

CDP & engagement (identity resolution, consented activation)

- **Twilio Segment** (global CDP), **CleverTap / MoEngage / WebEngage** (strong India presence for lifecycle engagement). (Recent MoEngage growth underscores category strength in India. Segment+3)

Social listening & app-store reputation

- **Sprinklr/Talkwalker** for social & public CX; **AppFollow** for app reviews. Sprinklr+2

Data catalog & governance

- **Atlan** (India-founded) for active metadata & AI governance. Atlan

  Notes: I'm vendor-neutral and unaffiliated. The list blends proven incumbents and India-born scale-ups; you only need a handful to start.

---

8) Policy, workflows, SOPs → policy-as-code

- **Codify SOPs** as rules (Rego/Cedar) with tests and approvals. Example (pseudo-Rego):

```arduino
allow_action["open_new_till"] {
  queue_p95_minutes > 7
  on_floor_staff_available >= 1
  sNPS_store_30m < target - 5
}
```

- **Change mgmt:** SOP PRs in Git; CI checks for fairness/privacy & simulation; rollout with feature flags; auto-revert on anomaly. (OPA/Cedar are designed for declarative, auditable policies. openpolicyagent.org+1)

---

9) Measurement & governance

- **Topline:** sNPS vs external NPS delta, NPS error (Brier/Calibration), NPS uplift from actions (causal).
- **Ops:** queue wait p95, OOS minutes, OTD%, FADR, app step-through rates, frustration sessions per 1k sessions. (Frustration sessions metric is standard in DXA. Fullstory)
- **Trust:** DPDP audits, consent rates, data minimization score, human-in-the-loop override logs. (DPDP assigns fiduciary duties and rights for principals. MeitY)

10) Risks & mitigations

- **Privacy & optics of cameras:** edge processing, no face ID, signage, aggregated outputs, retention caps; DPIA before rollout. (Axis/NVIDIA support on-edge analytics to reduce data risk. Axis Communications+1)
- **Silent NPS drift:** periodic calibration to survey; shadow mode for new models; experiment before wide rollout.
- **Action fatigue:** rate-limit tasks; weekly SOP retro to *remove* noisy triggers.

11) What this looks like day-to-day (once live)

- Store Manager opens the **Store Pulse**: "Queue p95 → 9.2 min (↑); Trigger fired; Till 4 opened; sNPS_store_30m improved +6."
- JioMart Product sees **Checkout Friction** spike; feature flag auto-rolled back; experiment records +1.1 pp conversion recovery; sNPS_journey↑.
- City Ops sees **ETA risk**; dispatch re-optimizes; proactive comms to customers; OTD preserved; sNPS_delivery steady.

12) A short, pointed Q&A list (optional)

1. **Your 12-layer names.** If you share them, I'll map one-to-one.
2. **Consent stance:** OK to use Wi-Fi pings? (Recommend *no*; stick to video edge + consented IDs per DPDP.) MeitY
3. **Current stack:** do you prefer Confluent/Snowflake/Databricks/AWS/GCP/Azure? (I'll tailor the blueprint.) Confluent+2
4. **Pilot picks:** 3 stores (formats), 1 category for shelf vision, 1 city & FC for quick-comm, and the "one killer" digital step (usually checkout).
5. **External survey partner:** Bain/Satmetrix/Medallia—any preference? (We'll need them for quarterly calibration.) Bain+1
6. **Workforce app:** existing frontline app to embed *Ops Copilot* or build greenfield?

Final take

This is a **doable** AI-first transformation: infer NPS from real-world behavior ("silent NPS"), act in seconds through policy-as-code SOPs, and prove causality with experiments—then validate quarterly with an external survey. Start small (3 stores, 1 city, 5% traffic), bias to edge privacy, and let the triggers teach you which SOPs to rewrite. The compounding effect is real: fewer queues and stockouts, faster deliveries, cleaner checkouts—and steadily rising promoter share.

If you'd like, I can turn this into a one-page RFP checklist for vendors and a 30-60-90 launch plan you can hand to your PMO next.

Yes need a targetted 60 day launch plan

Request for GPT-5 Pro
Reasoned for 4m 39s

Absolutely—here's a **targeted 60-day launch plan** to take "silent NPS" (no-survey inference) live across **3 pilot stores + 1 quick-commerce city + 5–15% of JioMart traffic**, wired into real-time actions, SOPs, and dashboards. It's organized as: what will be live by Day 60, the week-by-week plan, workstreams and deliverables, decision gates & acceptance criteria, and the exact playbooks, roles, and checklists you'll need.

By Day 60, you will have

Scope live

- **Stores:** 3 formats (e.g., hypermarket, supermarket, specialty).
- **Quick-commerce:** 1 city zone (1 MFC + ~10 service areas).
- **Digital:** 5% traffic in week 3 → 15% by week 6 on the checkout journey.

Outcomes live

- **Silent NPS v1 (sNPS)**: inferred promoter−detractor score, refreshed every 5 minutes at store/slot/journey level.
- **6 real-time action playbooks** auto-triggered from signals (queue, OOS, ETA risk, checkout friction, payment retries, negative-sentiment support).
- **Ops Copilot (MVP)** in frontline & city-ops app: tasks, SLAs, acks, escalations, and closed-loop outcomes.
- **Policy-as-code SOPs v1**: rules in Git with audit and simulation.
- **A/A and A/B experiments** running (queue & ETA playbooks).
- **Governance ready**: DPIA, signage pack, data minimization & retention, consent flows, model registry.

Hot path SLO

- Signal → action **≤ 5 seconds**; warm-path insights ≤ 60 minutes; daily cold-path reports.

60-day plan (week-by-week)

> **Critical path** runs through: (1) signal instrumentation, (2) streaming & feature store, (3) sNPS model v1, (4) action orchestration & Ops Copilot, (5) experiments & governance.

Week 1 — Mobilize, pick pilots, lock guardrails (Days 1–7)

- **Executive kickoff & goals** (2 hrs): confirm pilot scope, "no facial ID" stance, SLOs, KPIs.
- **Select pilots:** 3 stores (different formats), 1 quick-comm city, 1 digital step (checkout).
- **Vendor taps (lightweight)**: pick 1 for people counting/queue, 1 for digital experience analytics (DXA), keep in parallel with internal tools.
- **Governance/DPIA sprint 1:** data inventory, privacy-by-design controls, signage templates, retention policy (e.g., edge video frames discarded; metadata kept X days).
- **Architecture "thin slice"** confirmed: event bus, schema registry, stream processor, feature store, model registry, policy engine, orchestration, dashboards.
- **Team setup** (RACI, SLAs): name single-threaded owner per workstream; daily stand-up cadence.

Deliverables

- Pilot charter, SLOs, KPI tree, DPIA draft, system "thin-slice" diagram, backlogs per workstream.

Week 2 — Instrumentation & data contracts (Days 8–14)

- **Stores:** enable people counting & queue analytics on existing cameras; POS event taps (tender success/declines, item voids).
- **Quick-comm:** OMS/WMS/dispatch events (promised ETA, actual ETA, FADR, courier dwell).
- **Digital:** SDK tags for rage-clicks, errors, LCP/TTFB; error taxonomy; checkout event standards.
- **Data contracts**: schemas + PII minimization; employee-exclusion on counters; store zone map for dwell.
- **Streaming spine:** topics in event bus (Kafka or equivalent), schema registry; **dev → stage** pipelines.
- **Feature store bootstrap:** define 10 core features (see list below); online/offline parity.

Deliverables

- Instrumentation checklist (≥90% events emitting), 10 feature definitions, data contracts, topics & schemas, first dashboards (traffic, queues, checkout errors, ETA deltas).

Gate 1 (end of Week 2)

- **A/A test**: metrics stable across mirrors; **data freshness < 5s** for hot path; **privacy controls verified** (face blurring/employee exclusion on, retention timers).

Weeks 3–4 — sNPS v1 + policy-as-code + Ops Copilot MVP (Days 15–28)

- **sNPS v1**:
  - Labels from proxies: repeat purchase ≤14d, refund/complaint flags, successful substitutions, app ratings; small **external NPS sample** if available for calibration.
  - Train simple ensemble (logit + gradient boosting), **calibrate** (isotonic/Platt); compute **sNPS = mean(P(promoter)−P(detractor)) × 100** per entity (store/slot/journey).
  - **Drift monitors** (PSI) + daily calibration check.

- **Policy-as-code** (OPA/Cedar):
  - Author 6 rules (queue, OOS, ETA slip, checkout friction, payment retries, negative sentiment).
  - Unit tests + simulation; PR reviews by Ops + Privacy.
- **Ops Copilot (MVP)** within frontline app:
  - Task feed, SLA timers, "how-to" steps, ack/escalate/resolve; **action outcome capture** (e.g., till opened time, restock completed, call-back done).
- **Orchestration**: event-driven workflows (e.g., open till → notify trained associate → verify queue drop → auto-close or escalate).
- **Experimentation rig**: A/B framework live (can start with feature-flag tool + CUPED).

Deliverables

- sNPS v1 notebook + service endpoint, rules repo with tests, Ops Copilot build, orchestration flows, experiment playbook.

Gate 2 (end of Week 4)

- Shadow mode shows stable **sNPS coverage ≥ 80%** of interactions; rules fire in a sandbox; Ops Copilot end-to-end in **UAT**.

Week 5 — Controlled go-live & first A/Bs (Days 29–35)

- **Turn on** actions for pilot cohort:
  - **Queue playbook** (open till), **checkout friction rollback**, **ETA proactive comms**.
- **Start two experiments**:
  - **Store-level geo A/B** on queue playbook (treatment = policy on for 50% of time slots).
  - **City-level A/B** on ETA proactive comms.
- **On-floor training** (2× 30-min huddles): how tasks appear, SLA expectations, "fix-forward" etiquette for customers.
- **Ops control room** (lightweight): live wallboard, daily ops review.

Deliverables

- Experiment configs, training deck, wallboard, daily pilot report (conversion, queue p95, ETA deltas, actions per hour, action success rate).

Week 6 — Scale pilot & tighten loops (Days 36–42)

- **Expand digital** from 5% → **15% traffic**; expand store windows/slots.
- **Add 2 playbooks**: payment-retry assist; negative-sentiment callback.
- **SOP retro #1**: codify deltas that worked (e.g., queue staffing rule by hour); remove noisy triggers.
- **Model tweak**: add new features (payment retries, courier dwell variance).
- **Reliability hardening**: alerting on action lag >5s; auto-retries; dead-letter handling.

Deliverables

- SOP v1.1, model v1.1, SRE runbook, action reliability report.

Week 7 — Category & compliance depth (Days 43–49)

- **Shelf/OOS** for 1 category (computer vision or scan-based): gap detection → restock tasks.
- **DPIA sprint 2**: check live deployment, signage audit, retention audit, access logs; finalize model cards (intended use, limits).
- **CX root-cause**: tag top 3 drivers of detractors (from transcripts/reviews); open product eng tickets.

Deliverables

- OOS playbook live in 1 category; DPIA final; model cards; weekly driver report.

Week 8–9 (to Day 60) — Results, readiness, and next-scale (Days 50–60)

- **Experiment readouts** (90% power where possible): effect on conversion, queue p95, OTD; **delta on sNPS** vs control.
- **Playbook lifting plan**: which to scale to 10–15 stores & 2 cities in next 30 days.
- **Ops enablement**: finalize micro-trainings; publish "golden path" checklists.
- **Third-party survey prep** (kick off for Day 75–90) to calibrate sNPS; pre-stratify samples.

Deliverables

- Day-60 readout (exec deck), policy library v1.2, rollout plan (stores/cities/digital), backlog for next 90 days.

Workstreams, owners, and key deliverables

Workstream	DRI	Week-by-week deliverables
Edge Sensing (Store)	Store Tech Lead	W2: counters/queues live; W4: zone maps + calibrations; W7: OOS in 1 category
Digital Telemetry	Product Analytics Lead	W2: rage/error/perf events; W4: checkout friction detector; W5: rollback trigger
Quick-Comm Signals	City Ops Tech	W2: ETA & FADR; W5: ETA proactive comms A/B
Streaming & Features	Data Eng Lead	W2: topics/schemas/feature store; W3–4: online scoring path
Modeling & Causal	DS Lead	W3–4: sNPS v1 + calibration; W6: v1.1; W8: readouts
Policy-as-Code & Orchestration	Platform Lead	W3–4: rules + tests + workflows; W5: prod on; W6+: tuning
Ops Copilot & Training	Store Ops Lead	W4: MVP app; W5: go-live + trainings; W6–8: SOP retro
Governance (Privacy/AI Risk)	Privacy & Risk Lead	W1: DPIA draft; W2: controls; W7: DPIA final; continuous audits

Workstream	DRI	Week-by-week deliverables
Experimentation & Analytics	Experimentation PM	W4: rig; W5: A/Bs running; W8: results

Decision gates & acceptance criteria

Gate 1 (W2):

- Hot-path latency p95 ≤ 5s in stage; A/A variance within expected bounds; privacy controls verified; signage ready.
- **Go** to model build & policy coding.

Gate 2 (W4):

- sNPS coverage ≥ 80% of interactions; rule simulations pass; Ops Copilot UAT NPS (internal) ≥ 60.
- **Go** to controlled go-live & A/Bs.

Gate 3 (W6):

- Action reliability ≥ 99% (task created ≤ 5s, acked within SOP SLA), statistically valid early effects on operational KPIs.
- **Go** to scale pilot + add OOS playbook.

Day-60 Exit:

- ≥ 2 experiments with significant positive lift (or non-inferior with cost savings), sNPS tracks directionally with ops KPIs, governance complete, team ready for next-30 rollout.

The 6 launch playbooks (policy-as-code skeletons)

1. **Queue relief (store)**

```sql
when queue_p95_minutes > 7 AND on_floor_staff_available >= 1
then open_new_till WITH sla_minutes = 5
and confirm queue_p95_minutes < 5 within 15 minutes
```

2. **OOS restock (1 category to start)**

```sql
when shelf_gap_detected == true AND backroom_stock > 0
then create_restock_task WITH sla_minutes = 10
```

3. **Checkout friction (digital)**

```less
when checkout_error_rate_5m > 1.5% OR rage_click_pct > 3%
then rollback_feature_flag("checkout_variant") AND page_eng_alert
```

4. **ETA proactive comms (quick-comm)**

```csharp
when eta_slip_minutes >= 10 AND fadr_risk_high == true
then send_proactive_sms AND reroute_courier_if_possible
```

5. **Payment-retry assist**

```csharp
when payment_decline_rate_10m > threshold OR retries_per_user >= 2
then offer_alt_payment AND surface_oncall_chat
```

6. **Negative-sentiment callback (care)**

```sql
when call_sentiment <= -0.6 AND first_contact_resolution == false
then supervisor_callback WITH sla_minutes = 30
```

Core features (initial 10) for sNPS & triggers

1. **`queue_wait_p95_5m`**
2. **`dwell_time_ratio_zone`** (aisle dwell / store dwell)
3. **`shelf_gap_minutes`**
4. **`pos_decline_rate_5m`**
5. **`checkout_error_rate_5m`**
6. **`rage_click_pct_step`**
7. **`lcp_ms_p95_5m`**
8. **`eta_delta_minutes`** (actual − promised)
9. **`payment_retry_count_session`**
10. **`support_sentiment_score_1h`**

SOP changes you'll publish by Day 45

- **Queue staffing rule:** open additional till if p95 > 7 min for 10 minutes; greet waiting customers; close after p95 < 5 min for 15 minutes.

- **Restock rule (pilot category):** gap detected → restock within 10 minutes or log "no stock"; escalate to department lead.
- **ETA comms rule:** inform customer at +10 minutes slip; auto-credit goodwill within cap if >20 minutes.
- **Checkout rollback rule:** rollback on error/rage threshold breach; post-mortem within 24h.
- **Payment assist rule:** display alt payment, retry limit, route to chat after 2 retries.
- **Negative sentiment escalation:** supervisor callback within 30 minutes; document fix-forward steps.

All are encoded as **policy-as-code** with audit trails and versioned PRs.

Daily operating cadence (pilot)

- **08:30** Store Pulse: last 24h queues, OOS minutes, actions taken, sNPS trajectory.
- **09:30** City Ops: OTD, ETA slips, proactive comms outcomes.
- **10:30** Product/Engg: digital friction & rollbacks; bug triage.
- **16:30** Experiment check-in: guardrails, power/progress.
- **Weekly retro (Fri):** SOP patches, remove noisy triggers, add top-1 new trigger.

Team & resource snapshot (lean but real)

- **Program (1):** Overall PMO/Delivery lead.
- **Signals & Edge (3–4):** store video/IoT, quick-comm taps.
- **Data Platform (3–4):** streaming, feature store, lake/warehouse.
- **Modeling (2–3):** sNPS + causal/experiments.
- **Ops Copilot (3–4):** app, workflows, notifications.
- **Product Analytics (1–2):** digital telemetry & DXA.
- **Privacy & Risk (1–2):** DPIA, controls, audits.
- **Store/City Ops Champions (3):** per pilot site.

Risks to watch (with mitigations)

- **Privacy optics of cameras** → edge analytics only, signage, no facial ID, retention caps, employee exclusion.
- **Action fatigue** → rate-limit triggers, weekly retro, emphasize "fewer, higher-impact" tasks.
- **Model drift / mis-calibration** → small periodic external NPS sample; guardrails; shadow before deploy.
- **Data quality gaps** → A/A tests early; schema contracts; alert on null/volume anomalies.

What you'll measure during the 60 days

- **Reliability:** action creation ≤ 5s, failure rate < 1%, ack within SLA.
- **Ops KPIs:** queue p95, OOS minutes, OTD%, payment success rate, checkout completion rate.
- **Effect:** A/B delta on the above + **sNPS change**; correlation matrix for sanity.
- **Trust:** privacy control checks, policy audit logs, override counts.
- **Adoption:** Ops Copilot MAU, task completion %, time-to-ack, store manager feedback.

Quick appendices

A. Initial acceptance tests (sample)

- Trigger "queue" in a controlled window → new till staffed within 5 min → queue p95 drops below 5 min within 15 min in ≥ 70% cases.
- Checkout rollback fires within 2 minutes of breach; conversion recovers within 60 minutes vs control.
- ETA comms reduces cancellation rate by ≥ 10% in treatment vs control on slipped ETAs.

B. Minimal integration list

- POS events (tenders, voids) → stream;
- Camera counters/queues → edge → stream;
- OMS/WMS/dispatch → ETA events;
- App telemetry → SDK;
- Contact center → transcript + sentiment;
- Feature store + online inference; policy engine; workflow orchestrator; Ops Copilot app container.

C. Day-60 readout structure

- What we instrumented → reliability → actions → experiments → KPI/sNPS impact → SOP updates → risks → next-scale plan.

If you want, I can convert this into a one-page PMO Gantt and a policy-as-code starter repo (with the six rules, unit tests, and a simulation harness) you can hand to the teams.

Yes please

Request for GPT-5 Pro
Reasoned for 12m 0s
Done. I've created and packaged both items you asked for:

1) One-page PMO Gantt for the 60-day launch

- **PNG:** Download the Gantt (PNG)
- **PDF:** Download the Gantt (PDF)
- **Tasks table (CSV):** Download the task list

What's on the Gantt:

- Week-by-week bars (W1→W9) aligned to the plan we agreed.
- Decision gates clearly marked: **Gate 1 (Day 14), Gate 2 (Day 28), Gate 3 (Day 42), Day-60 Exit.**
- Owners shown on each bar (e.g., Program, Signals/Platform, DS/Platform/Ops, etc.).

  Tip: Drop the PNG into your weekly decks; keep the PDF as the PMO source of truth.

---

2) Policy-as-Code starter repo (OPA/Rego)

- **ZIP:** Download the starter repo

What's inside

```bash
policy-as-code-starter/
├── data/thresholds.json              # Tune limits & SLAs (e.g., queue p95=7m)
├── input-samples/                    # Ready-to-run sample inputs
│   ├── store_input_example.json      # queue_p95, staff_available…
│   ├── digital_input_example.json    # checkout_error_rate, rage_click_pct…
│   ├── delivery_input_example.json   # eta_slip_minutes, fadr_risk_high…
│   └── support_input_example.json    # call_sentiment, FCR…
├── policies/
│   ├── queue_relief.rego             # open_new_till
│   ├── oos_restock.rego              # create_restock_task
│   ├── checkout_friction.rego        # rollback_feature_flag
│   ├── eta_proactive_comms.rego      # proactive comms + goodwill flag
│   ├── payment_retry_assist.rego     # alt payment + chat assist
│   └── negative_sentiment_callback.rego# supervisor callback
├── tests/                            # Positive + negative test cases
│   ├── queue_relief_test.rego
│   ├── oos_restock_test.rego
│   ├── checkout_friction_test.rego
│   ├── eta_proactive_comms_test.rego
│   ├── payment_retry_assist_test.rego
│   └── negative_sentiment_callback_test.rego
├── scripts/simulate.sh               # One-command demo evals
├── Makefile                          # `make test`, `make eval-*`
└── .github/workflows/opa.yml         # CI: run OPA tests on every PR
```

How to run (quick start)

```bash
# 1) Install OPA (macOS/Linux/Windows): https://www.openpolicyagent.org/docs/latest/#running-opa
# 2) Unzip and cd into the repo
opa test -v policies tests data                # run the unit tests
opa eval -f pretty \
  -d policies -d data \
  -I input-samples/store_input_example.json \
  "data.policies.queue.actions"                # try a single policy
```

How to wire into your stack

- **Decision API pattern:**
  Your stream processor (Kafka/Flink) sends a compact event → a "Decision" service queries OPA (local or sidecar) for `data.policies.<policy>.actions` → if non-empty, orchestrator creates the task (Ops Copilot, OMS/WMS, notifications).
- **Config management:** Change thresholds in `data/thresholds.json`, commit, CI runs tests, and your CD rolls out the updated policies.
- **Auditability:** Persist (event_id, policy_version, actions[]) to your lakehouse/warehouse for drift and policy usage analytics.

Notes

- Everything intentionally avoids personal identifiers.
- Policies return **structured actions** with `type`, `reason`, and SLA/metadata—easy to convert into tasks.
- Unit tests cover both trigger and non-trigger scenarios for each rule.

Optional next step (fast)

If you'd like, I can also generate a **one-page vendor RFP checklist** and a **runbook** your PMO can hand to store and city ops for Day-1/Week-1.